

FUTON0

@multiarg futon0/devmap
 @title FUTON0 Development Map
 @audience futon-devs, reflective-agents, future-you
 @tone formal-analytic
 @style roadmap
 @allow-new-claims true
 @length any
 @factor Mindful apprehension (sati + sampajāñña)
 @IFR: FUTON0 is the mindfulness and futon-perception layer: the lived, experiential basis through which real-but-virtual futures (futons) become intelligible via symbolic mediation. It stabilises vitality, rhythm, clarity, and the conditions for "wondrous apprehension". FUTON0 provides a bio-mechanical interface to the extended cognitive architecture—the cybernetic floor from which the whole FUTON stack perceives, adapts, and co-evolves with its future.

@state Vitality telemetry is operational: the scanner, git activity pipeline, and Zoom R4 ingest run end-to-end, and the Stack HUD is the Futon0 surface (dedicated frame + context buffer) for vitality rows, boundary gaps, reminders, and voice-typing controls. Stack HUD snapshots are logged over time (daily JSONL plus compacted summaries). The data stores are still periodic JSON snapshots wired via cron/systemd rather than a real-time shared state.

@next Maintain and develop minimal models of vitality signals; cultivate periodic wondrous apprehension; and continue light futon-prototyping efforts with steady awareness, so that both clear devmaps and concrete implementations can emerge naturally.

! instantiated-by: Prototype 0 — System Vitality Indicators [?] / 系统 [?] / 市]

+context: FUTON0 provides the minimal embodied/creative signals required for early detection of imbalance, while avoiding expansion of lifelogging overhead.
 +if: You want dependable early warning of drift in health, energy, or expressive rhythm.
 +however: Existing signals are scattered (notes, reminders, implicit habits) and therefore unreliable as a coherent early-warning interface.
 +then: Use only "existing traceable activity" as indicators: Tatami sessions, git commits, filesystem timestamps, and micro-recording cadence. Add one lightweight prompt-based behaviour: a weekly tai-chi intention check surfaced automatically by the FUTON3 HUD. No manual lifelogging is introduced.
 +because: Indicators remain sustainable only when they create no friction. A minimal interface built from pre-existing behavioural traces provides consistent, low-noise signals that prevent hidden imbalance from migrating up-stack.

+ status[done]
 +evidence:
 - The FUTON0 clock already exists: Tatami/API chat sessions automatically emit timestamped events, providing a ground truth for daily activity.
 - Git commits, recording timestamps, and filesystem modification times supply ambient movement/engagement signals with zero added logging.
 - FUTON1 work-session tags already show >90% end-of-session coverage.
 +next: Build the cron-based vitality scanner; implement the weekly tai-chi HUD prompt; run a 30–45 day pilot to stabilise thresholds; and refine the vitality interface only after reviewing correlations with mood, focus, and WIP drift.
 +next-evidence:
 - Extend the cron scanner to generate a daily summary of filesystem activity, recording cadence, and Tatami timestamps.
 - Capture tai-chi intention Y/N responses in FUTON3, creating a new sparse evidentiary channel without manual logging.
 - Implement a lightweight inference pass

FUTON1

@multiarg futon1/devmap
 @title FUTON1 Development Map
 @audience futon-devs, infra-agents, future-you
 @tone formal-analytic
 @style roadmap
 @allow-new-claims true
 @length any
 @factor Keen investigation (dhammavacaya)
 @IFR:
FUTON1 provides a deterministic, storage-oriented substrate that maintains a coherent graph of facts across sessions. It keeps Datascript (fast) and XTDB (persistent) in reproducible alignment so that the same world state appears reliably across restarts and environments. It exposes a minimal CLI and HTTP surface for ingesting text, inspecting entities and relations, and emitting stable summaries—such as focus headers—that downstream code can depend on without embedding persistence or orchestration logic.

FUTON1 supports keen investigation by making the fact graph inspectable and faithful to stored state. Its role is limited to storage, hydration, and deterministic summarisation: it does not prescribe inference or policy. When this layer operates cleanly, other futons inherit a durable and predictable world model on which more complex behaviours can be built.

@state The deterministic demo/client workflow and Datascript+XTDB mirroring are operational, with sessions hydrating from XT on boot and falling back to the legacy log + snapshot when XTDB is disabled. Futon1 remains storage- and durability-focused: 'apps/demo' is the interactive entry point and the HTTP API module ('clojure-M/api' / 'apps/api') is the external client surface.

Operational caveats: multiple Futon processes against the same 'BASIC_CHAT_DATA_DIR' can trigger RocksDB lock errors, and open-world ingest treats questions as assertions (there is no question→Datalog mapping yet).
 @next
 Consolidate the documentation and contracts that define how text becomes graph state; stabilise the naming and configuration surfaces shared across CLI/API, and ingest workflows; and refine test coverage so the deterministic substrate remains reliable as modules evolve. Maintain FUTON1's minimal scope—storage, hydration, and inspection—while ensuring that higher futons can depend on its interfaces without ambiguity.

! instantiated-by: Prototype 0 — Baseline Deterministic Substrate [?] / 四 [?] / 基

+context: FUTON1 provides a deterministic storage substrate built from the demo/client workflow, the NLP interface, graph-memory, and XTDB persistence. These components together define the minimal end-to-end path by which text becomes durable graph state and can be inspected through the CLI or HTTP API.
 +if: You want a stable understanding of what FUTON1 currently guarantees—its ingest pipeline, hydration behaviour, and contract surfaces—before extending or modifying any part of the stack.
 +however: Documentation and examples are distributed across several READMEs; the overall "stack story" from ingest to XTDB is implicit rather than expressed as a single reference artifact, and the exact boundaries of the deterministic substrate are easy to lose track of during development.
 +then: Capture a baseline snapshot of the current system: record a concise architecture note describing the pipeline (ingest → NLP → graph-memory → XTDB), freeze representative transcripts from the demo/client workflow, and tag a reference point that reflects the present deterministic behaviour.
 +because: A clear baseline makes the substrate's guarantees explicit, helps detect regressions in hydration or pipeline flow, and provides a stable reference against which future refinements or

FUTON2

@multiarg futon2/devmap
 @title FUTON2 Development Map
 @audience futon-devs, simulation-agents, future-you
 @tone formal-analytic
 @style roadmap
 @allow-new-claims true
 @length any
 @factor Keen investigation (dhammavacaya)
 @IFR:
FUTON2 is the viriya (energy) chamber: a narratable Active Inference engine whose ants are first demonstrators for AIF agent species that will eventually crawl FUTON1 graph memory, export state to higher futons, and turn insight into kinetic practice across the network.

@state Ant sim + AIF loop run end-to-end with analyzer hooks, and cyber-ants can swap in for the classic faction via Futon3-derived patterns. Architecture baseline documented in stack-baseline.md. Observation property tests strong (440+ iterations). Hunger/tau goldens recorded (6 scenarios). Policy goldens recorded (5 scenarios). Gaps: formal semantics docs and white-space tests missing.

! instantiated-by: Prototype 0 — AIF Stack Baseline (observe → perceive → affect → core) [?] / 四 [?] / 基

+context: FUTON2 runs a full Active Inference loop in the ant war-game, wiring 'ants.aif.observe', 'ants.aif.perceive', 'ants.aif.affect', , and 'ants.aif.core' into one tick-by-tick pipeline.
 +if: You want a stable story that shows how observation, prediction errors, precision, and actions integrate per tick.
 +however: Modules are individually solid but no single document freezes hunger/tau semantics or traces the whole stack.

+then: Capture a top-level architecture note from sensory keys through world updates, lock current hunger and predictive micro-step parameters, and add four golden AIF microtraces.
 +because: A documented baseline prevents accidental drift before you optimise policy, analyzers, or mechanics.

+evidence:
 - 'futon2/doc/stack-baseline.md' (76 lines): Documents 5-stage pipeline (observe→perceive→affect→policy→act), locks baseline parameters (max-steps=4, alpha=0.45, beta=0.28), lists observation vector ABI.

- 'futon2/src/ants/aif/core.clj' (209 lines): 'aif-step()' orchestrates full tick, 'default-aif-config' defines all nested config (preferences, precision, actions, efe, modes).

- 'futon2/src/ants/aif/observe.clj' (185 lines): 'g-observe()' normalizes 14 channels to [0,1].

- 'futon2/src/ants/aif/perceive.clj' (165 lines): Multi-step predictive coding with trace output.

- 'futon2/src/ants/aif/affect.clj' (240 lines): 'tick-hunger()', 'hunger->tau()', 'modulate-precsions()', 'update-tau()' .

- 'futon2/src/ants/aif/policy.clj' (841 lines): EFE-based action selection with mode-conditioned biases.

- 'futon2/doc/trace/single_tick.edn' + 'futon2/doc/trace/single_tick.svg' : Deterministic trace artifact + plot.

- 'futon2/test/resources/goldens/micro-traces.edn' : 4 golden perception traces (baseline, hungry-needy, sated-on-trail, homebound-low-phr).

+next-evidence:

- 'doc/trace-semantics.md' explaining micro-step semantics and why errors are 0.0 in goldens.

- 'doc/hunger-tau-coupling.md' documenting when update-tau fires vs perceive-internal hunger updates.

- 4 complete tick traces with commentary (forage, return, explore, defend scenarios).

! instantiated-by: Prototype 1 — Observation Layer Hardening [?] / 四 [?] / 四

+context: 'g-serve()' normalises sensory evidence (food, pheromones, proximity, novelty, reserves) into the vectors passed downstream.
 +if: You want deterministic observations across

FUTON3

@multiarg futon3/devmap
 @title FUTON3 Development Map
 @audience futon-devs, sandbox-agents, future-you
 @tone formal-analytic
 @style roadmap
 @allow-new-claims true
 @length any
 @factor Energy (viriya)
 @IFR:
FUTON3 is the place where the system turns messy activity into organised knowledge. It collects what happens, checks it against a library of shared patterns, and produces clear, trustworthy records that other futons can use to learn, improve, and act. FUTON3 makes everyday work auditable, meaningful, and easier to build on.

@state Transport runs reliably, events are logged, and patterns can be checked.

Tatami sessions already feed structured evidence downstream, but that evidence does not yet flow back into pattern updates or storage-backed history. We have an early pattern catalog, working embeddings, and a first version of the workday pipeline. What remains is tightening those links – coupling Tatami evidence into pattern updates, tying checks into storage, producing clearer histories, and building the stable interfaces that other futons can depend on for coordinated reasoning.

! instantiated-by: Prototype 0 — Running Since 29 July 2025 [?] / 久 [?] / 用

+context: Prototype 0 establishes the symbolic foundations, artefact inventory, and DREAM logic.
 +if: You want a coherent first cycle of a year-long 12-prototype run.

+however: Symbolic activation does not yet form a concentrated practice ecology.
 +then: Define FUTON3 as the layer where prototypes from F4 are refined into systems—pattern libraries, CT riffs, musical grammars, improvisation logics, Nonstarter economics—so concentrated practice can run independent of outcome-chasing.

+because: Without a stable concentrator layer, FUTON4's tranquility and FUTON6's argument power cannot synergise.

+status[ready]

+evidence:
 - 'spine.org' — Master orchestration file with tangling machinery (lines 109–103).
 - 'dev/bootstrap.el' — Repeatable build harness for interactive + batch builds.

+if: You want an execution substrate where pattern obligations can be posed, evaluated, and recorded with deterministic transcripts.

+however: The baseline hello/event/session flows are documented as "non-essential documentation" instead of "proofwork IO contracts."

+evidence:

- 'test/arxana-bootstrap-test.el' — Baseline harness validation (3 tests).

- 'arxana-tangled.el' — Generated single-source output (3752 lines) proving round-trip.

! instantiated-by: Prototype 1 — Patterns of Improvisation & Meta-Pattern Formalism [?] / 五 [?] / 工

+context: FUTON3 still rides the MUSN transport (WS+HTTP bus, SAFE/ADMIN REPL, NDJSON ingest), but its primary job is to host pattern checks and informal proof states, not free-roaming agents.

+if: You want an execution substrate where pattern obligations can be posed, evaluated, and recorded with deterministic transcripts.

+however: The baseline hello/event/session flows are documented as "non-essential documentation" instead of "proofwork IO contracts."

+evidence:

- 'agents.md' — Development manifesto describing Phase 1/2 strategy.

- 'docs/reanimation-plan.org' — M0–M6 milestone tracker with checkbox progress.

+however: The baseline hello/event/session flows are documented as "non-essential documentation" instead of "proofwork IO contracts."

+evidence:

- 'test/arxana-bootstrapping-test.el' — Baseline harness validation (3 tests).

- 'arxana-tangled.el' — Generated single-source output (3752 lines) proving round-trip.

! instantiated-by: Prototype 1 — Arxana Reboot (Week-long target) [?] / 久 [?] / 用

+context: Legacy literate-programming structure exists but is dormant.

+if: You want stable parsing, node navigation, and save-back-to-LaTeX working.

+however: Legacy Emacs-Lisp and refactored Clojure schemas diverge.

+evidence:

- 'docs/protocol.md', 'src/f2/router.clj', 'test/f2/router-test.clj' (100-event idempotency + export/run/status acceptance), and 'test/transport-test.clj' (per-client disconnect logging).

+next-evidence:

- :type Spec.id "transport-contract-check-jobs-v1" :description "Define a frozen transport contract for pattern-apply, gap-report, and trail-capture jobs including message shapes, sequencing, and log guarantees" :acceptance-criteria ["Spec enumerates request/response schemas, NDJSON log records, and required invariants for all three check jobs."] :fails if there is no reliable A->B association between proposal and evidence collection; must be machine-checkable now and upgradeable for stronger logics later.] :suggested-location "docs/protocol/transport-contract-v1.md" :timebox "2026-01-05"]

- :type Spec.id "golden-transcripts-check-jobs" :description "Specify canonical golden transcripts for pattern-apply, gap-report, and trail-capture runs that deterministically exercise the frozen transport contract." :acceptance-criteria [" Each transcript is reproducible from the sandbox profile and matches byte-for-byte expected NDJSON output."] :fails if there is no reliable A->B association between proposal and evidence collection; must be machine-checkable now and upgradeable for stronger logics later.] :suggested-location "docs/protocol/golden-transcripts.md" :timebox "2026-01-05"]

+evidence:

- 'arxana-adjacency-test.el', 'arxana-saving-test.el', 'arxana-store-test.el' .

+Storage bridge complete: 'dev/arxana-store.el'

(350+ lines) with 12 Futon API helpers,

documented in 'docs/storage-bridge.org'

with curl examples.

- Article lifecycle: 'dev/arxana-article.el'

(metadata sync, deletion propagation).

- Scholia authoring: 'dev/arxana-scholium.el'

(compose, region capture, authoring mode).

- Derivation previews: 'dev/arxana-derivation.el'

(collapsible previews, color overlays,

toggle commands for inclusion/transclusion/iden-

fication).

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is* the specification that generates implementation.

+because: These meta-patterns form the FUTONS "RNA pack": a finite alphabet of transformations grounded in CA control features, with category theory providing the composition rules.

The formal pipeline: **patterns → category-theoretic schema → code generation** (via LLM or other methods). This enables executable documentation—the pattern composition *is*

(FUTON3 or FUTON4) to interpret these signals for trend detection and early-warning triggers.

! instantiated-by: Prototype 1 — Epistemic Rhythm & Salients [心/弓/桥]

+context: FUTON3 stabilises your epistemic rhythm: a weekly/multi-weekly cycle of review, planning, and salience-tracking that keeps work across all spheres coherent while avoiding data fragmentation or over-tracking.
+if: You want to programme yourself the way you programme systems—maintaining legible salients across institutional, consulting, technical, reflective, and infrastructural spheres with minimal friction.
+however: Without a normalised rhythm, weekly reviews slip; projects drift outside the canonical Org structure; and WIP silently grows past safe limits.
+then: Hold a weekly AOB/Ambitions sync; treat Org as the canonical plan/export store; mirror all active projects into a minimal FUTON1 substrate; review sphere balance weekly; and enforce the WIP cap (>3 active tracks requires pausing one). Cron/HUD integration surfaces upcoming reviews (“Weekly review in 3 days,” “Monthly audit approaching”) without additional logging.
+because: A stable epistemic rhythm prevents unbounded cognitive load “and makes life feel lighter, more joyful, and more tractable”. When salients reappear on schedule and WIP remains bounded, drift does not accumulate and affect settles.

+evidence:
- Org already represents the canonical project hierarchy (AOB, Ambitions).
- Weekly AOB/Ambitions sessions are timestamped (Tatami/API or calendar).
- FUTON1 tags provide >90% work-session closure coverage.
- Sphere balance and WIP creep can already be inferred from Org + FUTON1.
- status[ready-pending-final-decisions-and-testing]

- ‘futon0/futon0_rhythm.md’ (cadence, envelope logging, buildout plan).
- ‘futon0/futon0_protocol.md’ (open decisions + affect sourcing).
- ‘futon0/data/backup_cadence.json’, ‘futon0/data/backup_report_example.json’, ‘futon0/data/backup_experiments.json’, ‘futon0/data/affect_markers.json’.
- ‘futon0/scripts/futon0/backup_snapshot.clj’, ‘futon0/scripts/futon0/rhythm/envelope.clj’, ‘futon0/scripts/futon0/rhythm/salients.clj’, ‘futon0/scripts/futon0/rhythm/experiments.clj’, ‘futon0/scripts/futon0/rhythm/quarterly.clj’, ‘futon0/scripts/futon0/rhythm/dry_run.clj’.
+next: Consolidate the weekly sync into a fixed time; prune stale Org headings monthly; ensure every active project has an Org headline and FUTON1 tag; integrate HUD reminders for weekly/monthly cadence; and refine the WIP audit to produce actionable “pause/don’t-start” prompts.

+next-evidence:
- Cron job generates weekly/mid-week/monthly reminders surfaced in the HUD.
- FUTON3 salience panel merges Org + FUTON1 metadata into a weekly salients snapshot with diff from prior weeks.
- Quarterly rhythm analysis correlates stable cadence with subjective affect measures (energy, focus, lightness) to test the experiential claim.

! instantiated-by: Prototype 2 — Core Obligations Contract [业/业/书]

@multiaugfuton0/prototype2
@title Core Obligations Contract
@audience futon-devs, reflective-agents
@tone analytic-concise
@style flexiarg
+context: FUTON0 must uphold the fundamental constraints that keep the entire stack viable: sustained alignment with your current role, movement toward a more suitably aligned next role, and protection of basic health baselines. This layer also ensures that Prototype 1’s analytics loop remains genuinely useful rather than self-perpetuating.

+if: You want your working life to remain healthy and coherent while moving toward a better-aligned future role, without sacrificing present obligations or drifting into avoidant behaviour.
+however: When attention shifts to higher layers,

extensions can be evaluated.

+evidence:
- README documents deterministic demo/client behaviour, NLP interface, graph-memory schema, and XTDB hydration.
- ‘apps/demo’, ‘apps/client’, ‘apps/nlp-interface’, ‘apps/graph-memory’, and ‘apps/open-world-ingest’ form the operational substrate.
- XTDB inspection and hydration behaviour are reproducible through existing scripts and aliases.
- ‘/api/alpha/meta/model/registry’ exposes live model descriptors for docbook, open-world ingest, and the meta-model, with ‘/api/alpha/meta/model/queue’ reporting pending coverage.
- Open-world invariants pass end-to-end, including stub detection and relation type registration.
+next: Compile the architecture note; select and store a minimal set of canonical transcripts; and publish them alongside a tagged release that marks the baseline behaviour.
- Add a lightweight ‘/api/alpha/entities/latest?type=...’ endpoint (or equivalent) so downstream tools can fetch only the most recent ‘clock-out/summary’ (or other entity type) instead of paging the entire history when verifying persistence.
+next-evidence:
- A dated, versioned snapshot containing architecture notes and representative transcripts.
- CI or local reproducibility checks demonstrating identical XTDB and focus-header outputs under the baseline tag.
- A single README-level reference pointing to the baseline for future comparison.

! instantiated-by: Prototype 1 — Demo/Client Determinism & Focus-Header Instrumentation [计/每/口/律]

+context: The ‘apps/demo’ and ‘apps/client’ workflow is the canonical entry point for interrogating the stack. It emits deterministic focus headers, hydrates from XTDB on startup, and exercises the NLP interface and graph-memory module together, exposing how ingest, mutation, and inspection behave in practice.
+if: You want the CLI and client workflow to act as reliable instrumentation for investigating hydration, salience updates, and focus-header formation—producing traces that can be compared across runs and environments.
+however: Existing smoke tests cover only narrow paths, and the behaviour of hydration, persistence, and focus-header emission under varied configurations is not yet exercised systematically. Legacy assumptions from earlier protocols still colour expectations about determinism.
+then: Establish a minimal determinism suite for the demo/client workflow: select a small set of representative sessions, fix their configuration, and record the expected XTDB and focus-header outputs. Treat these as canonical fixtures checked under CI so that changes in hydration, pipeline flow, or header structure become immediately visible.

+because: Treating the demo/client workflow as calibrated instrumentation—rather than an informal demo—provides a stable, reproducible surface for higher futons, lets developers detect drift in persistence or header logic, and clarifies exactly what guarantees FUTON1 provides.

+evidence:
- README documents deterministic demo/client behaviour, focus-header output, and XTDB hydration.

- Golden fixtures already exist for graph-memory and NLP layers.

- XT inspection paths (‘M:storage/inspect’) demonstrate reproducible hydration and salience persistence.

- CI green on commit 761ea2 for ‘Futon1 Deterministic Stack’ (apps/client, apps/graph-memory, apps/nlp-interface).

+next: Curate a minimal determinism suite; freeze representative focus-header and XTDB outputs; and integrate these checks into CI as regression indicators.

+next-evidence:
- A small directory of canonical demo/client transcripts with corresponding expected XTDB states.

seeds and a nailed-down vector ABI for later ML. +however: Normalisation bounds and neighbour logic assumptions are scattered across helpers. +then: Write property tests that enforce 0-1 ranges, document the ‘sense->vector’ ordering, and add edge-case tests for grid borders, max-distance tweaks, and degenerate pheromone fields.

+because: Reliable observations keep perception and policy stable no matter the world configuration.
+evidence:

- ‘futon2/test/ants/aif/observe_test.clj’ (207 lines): Property tests + unit tests.
- ‘g-observe-values-clamped’: 200 random worlds, all 16 keys verified ∈ [0,1].
- ‘g-observe-border-neighbor-means’: 120 iterations testing edge/corner calculations.
- ‘degenerate-field-normalizes-to-zero’: Handles max-food gracefully.
- ‘sense-vector-ordering’: Locks 14-element vector ABI.
- ‘futon2/src/ants/aif/observe.clj’: ‘clamp01()’, ‘normalize()’ enforce bounds; all channels pass through normalization before return (lines 161-177).
- ‘futon2/doc/stack-baseline.md’ (lines 7-9): Lists observation ABI ordering.
- ‘futon2/README.md’ (lines 88-93): Documents sense->vector keys (needs update: lists 13, actual is 14).
+next-evidence:

- ‘doc/observation-spec.md’ listing all normalization bounds (max-food, max-pher, max-dist, max-reserve).
- White-space detection property tests (‘white’ field currently untested).
- Documentation of 8-neighborhood offset logic.

! instantiated-by: Prototype 2 — Predictive Coding Microcycle Clarity [虫/乡/虫/叉]

+context: Perception performs multi-step updates of mu, tau, Pi-o, goals, and weighted errors with annealing.
+if: You want inspectors (or future Arxana bridges) to understand belief updates.
+however: Hunger updates, sensory prediction, and error accumulation are interwoven without a canonical trace.

+then: Publish a single-tick perception trace with plots of hunger, tau, and error per micro-step, and formalise an interface for extracting a mu/precision snapshot.

+because: Clear cycles make later integrations with futon0 symbolic memory far easier.

+evidence:

- ‘test/resources/goldens/microtraces.edn’ (24 lines): 4 golden scenarios with per-step traces.
- ‘baseline’: h 0.66→0.79, tau 1.47→0.81 over 4 steps.

- ‘hungry-near-enemy’: h jumps to 1.0, tau floors.

- ‘sated-on-trail’: h stays low 0.22→0.33, tau peaks 2.06→1.5.

- ‘homebound-low-pher’: h 0.55→1.0, tau drops 1.76→0.33.

- ‘test/ants/aif/microtrace_regression_test.clj’ (80 lines): Runs 4 scenarios through ‘aif-step’, compares ‘perception:trace’ to golden, locks baseline parameters.

- ‘tools/trace/tick.clj’ (151 lines): Generates ‘doc/trace/single_tick.edn’ + SVG plot.

- ‘doc/trace/single_tick.svg’: Renders tau (red), h (blue), error (purple) with legend.

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘doc/trace-semantics.md’ formalizing micro-step semantics (why errors are 0.0, annealing formula).

- Goal-bias mechanism documented and tested (blending toward home/enemy based on cargo).

! instantiated-by: Prototype 3 — Hunger & Precision Dynamics Goldens [火/六/眼]

+context: Functions such as ‘tick-hunger’, ‘update-hunger’, ‘update-tau’, and ‘modulate-precisions’ control the most sensitive dynamics.

+if: You want reproducible affective behaviour that survives refactors.

+however: There are no golden tests for tau evolution, hunger saturation, or reserve-driven adjustments.

+next-evidence:
- A small directory of canonical demo/client transcripts with corresponding expected XTDB states.

runtime” :description “ Document the sandbox profile as the canonical proofwork runtime and explain how to execute check jobs from the Clojure runner.” :acceptance-criteria “[Spec explains sandbox configuration, versioning rules, and step-by-step runner commands tied to transcript generation.” ” Fails if there is no reliable A>B association between proposal and evidence collection (i.e., we need a clear and deterministic correspondence between antecedents A and measurable evidence B); must be machine-checkable now and upgradeable for stronger logics later.”] :suggested-location “ docs/sandbox/README.md” :timebox “ 2026-01-05” }

! instantiated-by: Prototype 1 — Pattern Canon & Standard Library [目/目]

+context: A small flexiarg/pattern library already exists (Nonstarter, Proto/Allow-Works, etc.) but is scattered between files.

+if: You want FUTON3 to be the place where patterns are curated with metadata, linking obligations, and proof status.

+however: There is no standardized repository or schema for “pattern definitions + applicability notes + links to Futon devmaps.”

+then: Stand up a ‘f3.patterns’ store (ID, clauses, references, pāramitā tags, fruits/orbs weights) populated with the current standard library and backfilled with source links. The pattern store should be addressable/interactable from across the stack. So, it should use futon1 storage, with contents feature comple with the filesystem devmaps and library, which provide a “checked out working copy”; the store needs to be easy to maintain relative to these checked out copies. Patterns managed by the store should be inspectable, linkable, and editable with Arxana (futon4), and subject to ongoing improvement via aob-chatgpt(futon3).

+next-evidence:

- ‘docs/migration-guide.org’ documenting import paths for historical .epub/.tex formats.

- ‘github/workflows/test.yml’ or equivalent CI configuration.

- Git tag marking the “revived” Prototype 1 release.

! instantiated-by: Prototype 2 — Graph Memory (Datascript) [目/久/手/用]

+context: The hyperedge schema for Arxana exists on paper but not in running code.

+if: You want argument fragments to become addressable nodes with provenance.

+however: there is no minimal memory manager wiring Arxana buffers to Datascript.

+then: Implement ‘add-node’, ‘activate-node’, and ‘link-node’; persist them via Datascript; replay a session to prove round-trips.

+because: without a living graph memory, Arxana never escapes editor cosplay.

+evidence:

- We have an a basic working version of the store in futon1, including ingest, ‘resources/pattern_store.edn’, ‘src/futon3/pattern_store.clj’, and ‘test/pattern_store_test.clj’ keep the catalog synced with live tests + README references.

+next-evidence:

- {type :Golden id “ pattern-store-loads-and-enumerates” :description “ A golden test proves the f3.patterns store can be loaded at runtime and enumerates all canonical pattern IDs.” :acceptance-criteria “[Loading the pattern store via the public Futon3 API returns a non-empty, stable set of pattern IDs matching the standard library inventory.” ” Fails if the canonical store only exists on disk and is not accessible to the stack.”] :suggested-location “ test/futon3/pattern_store_golden_test.clj” :timebox “ 2026-01-06” }

- {type :Golden id “ pattern-metadata-and-links-complete” :description “ A golden test asserts that each canonical pattern includes required metadata fields and valid source or devmap links.” :acceptance-criteria “[Every pattern entry exposes ID, clauses, references, pāramitā tags, fruits/orbs weights, and at least one resolvable source link.” ” Fails if the canonical store only exists on disk and is not accessible to the stack.”] :suggested-location “ test/futon3/pattern_metadata_golden_test.clj” :timebox “ 2026-01-06” }

- {type :Golden id “ pattern-store-matches-standard-library” :description “ A golden test verifies that all patterns from the existing standard library are present in f3.patterns with no omissions.” :acceptance-criteria “[The set of patterns in f3.patterns exactly matches the known Nonstarter and Proto/Allow-Works library with no missing or extra entries.” ” Fails if the canonical store only exists on disk and is not accessible to the stack.”] :suggested-location “ test/futon3/pattern_library_sync_golden_test.clj” :timebox “ 2026-01-06” }

- {type :Golden id “ pattern-store-matches-standard-library” :description “ A golden test verifies that all patterns from the existing standard library are present in f3.patterns with no omissions.” :acceptance-criteria “[The set of patterns in f3.patterns exactly matches the known Nonstarter and Proto/Allow-Works library with no missing or extra entries.” ” Fails if the canonical store only exists on disk and is not accessible to the stack.”] :suggested-location “ test/futon3/pattern_library_sync_golden_test.clj” :timebox “ 2026-01-06” }

- {type :Golden id “ pattern-store-matches-standard-library” :description “ A golden test verifies that all patterns from the existing standard library are present in f3.patterns with no omissions.” :acceptance-criteria “[The set of patterns in f3.patterns exactly matches the known Nonstarter and Proto/Allow-Works library with no missing or extra entries.” ” Fails if the canonical store only exists on disk and is not accessible to the stack.”] :suggested-location “ test/futon3/pattern_library_sync_golden_test.clj” :timebox “ 2026-01-06” }

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

- ‘src/ants/aif/perceive.clj’ (lines 133-164): Loop outputs ‘{:tau,:h,:error}’ per micro-step.

+next-evidence:

core obligations can slide into the background; health erosion becomes invisible; and epistemic-review routines can quietly expand beyond their real value.

+then: Maintain consistently and measurably strong alignment with the current role while increasing alignment with a future role; keep health baselines intact; and periodically check whether Prototype 1's analytics loop is delivering practical benefit.

+because: Upholding alignment and constraints at the base prevents the entire FUTON stack from drifting into incoherence or overload. When obligations, health, and review practices stay grounded, forward movement becomes stable and real.

+evidence:

- Health baselines visible through Prototype 0 (movement/sleep Y/N).
- Current-role performance inferred from deliverables, responsiveness, and deadline stability.

- FUTON1 tags showing adequate work-session closure coverage.

+next: Make alignment in the current role more explicit (quick weekly pulse-check); establish a small, low-friction trajectory toward the next aligned role; and add a brief weekly check on the usefulness of Prototype 1.

+next-evidence:

- Cron-backed extraction of role-search traces (calendar items, files, emails) to show gentle forward movement.

- A FUTON3 prompt asking "Did this week's review help maintain alignment?" as a simple measure of analytics-loop value.

- Correlations between Prototype 0 vitality signals and stable current-role performance to validate the alignment hypothesis.

! instantiated-by: Prototype 3 — Ambition 0 Meta-Managed Infrastructure [💡/叉 ❤️/稳]

@multiarg futon0/prototype3
@title Ambition 0 — Meta-Managed Infrastructure
@audience futon-devs, reflective-agents
@tone analytic-concise
@style flexiarg

+context: FUTON0 needs an infrastructure layer that not only keeps reality and representation aligned, but also increases the system's capacity for alignment over time (Ashby-style requisite variety).

+if: You want your tools to help you notice, raise, and refine signals that support increasing alignment in work, health, and direction.

+however: Interfaces drift, signals blur, and representational structures lose fidelity. When this happens, the system cannot increase alignment because it cannot perceive what is changing or what needs attention.

+then: Treat FUTON1 as the source of truth; use Org as a lightweight interface to it; prune dashboards to highlight only rising or fading salients; and run a brief reconciliation loop that raises proto-pattern signals for review.

+because: This layer provides the Markov blanket through which alignment can improve. Prototype 3 acts as the experience-design prototype of FUTON3: it raises signals, surfaces proto-patterns, and increases representational variety, while FUTON3 will later formalise this as proof-theoretic salience certification.

+evidence:

- Org and FUTON1 currently reflect overlapping state that allows detection of divergence.

- Dashboards already reveal what is becoming noisy or inert.

- Backups and timestamps show which elements are live or stale.

+next: Migrate remaining project state into FUTON1; reduce Org to a thin interface; and introduce a weekly "rising/fading signals" pane in FUTON3 to surface proto-patterns for reflection.

+next-evidence:

- Divergence checks that quantify where representation is lagging reality.

- Signal-rise/signal-fade detectors for dashboards and headings.

- FUTON3-generated proto-pattern stubs based on behavioural traces.

! instantiated-by: Prototype 4 — Ambition 1 Hyperreal Enterprises Practice [💡/习 📈/已]

@multiarg futon0/prototype4

@title Ambition 1 — Hyperreal Enterprises Practice

- CI reports confirming stability of determinism across tagged commits.

- Clear README references linking these fixtures to FUTON1's guarantees.

! instantiated-by: Prototype 2 — NLP Interface Consolidation [💡/申 🎯/内]

+context: 'apps/nlp-interface' provides deterministic NER/POS tagging and a small set of pattern-recognition hooks used across ingest and demo/client workflows. It is the shared entry point for turning raw text into structured tokens that the rest of FUTON1 can rely on without ambiguity.

+if: You want a single, well-defined NLP module whose behaviour is testable, documented, and independent of the surrounding CLI or transport logic.

+however: The current README gives only a minimal overview of the pipeline, and the boundaries between tokenisation, tagging, chunking, and downstream pattern recognition remain implicit in code rather than expressed as a clear API. Tests cover specific fixtures but not the full staged behaviour.

+then: Make the pipeline stages explicit: document and expose the tokenise → tag → chunk → recognise flow, write unit tests for each stage, and provide a brief, focused README describing the module's public API, expected inputs and outputs, and how upstream components should depend on it.

+because: A consolidated and documented NLP interface provides a stable grounding for ingest and demo/client workflows, reduces accidental coupling, and makes it easier to extend or revise individual stages without destabilising the rest of FUTON1.

+evidence:

- README describes deterministic NER/POS behaviour and shared pipeline usage.
- Test runners exist for NLP fixtures ('apps/nlp-interface/test').

- Pattern-recognition hooks are already used in ingest and CLI flows.

+next: Clarify the staged pipeline, expand tests to cover each stage, and produce a concise README that reflects the actual API and expected data shapes.

+next-evidence:

- Stage-by-stage unit tests passing under CI.
- Updated README documenting pipeline flow and public functions.

- Agreement across ingest and demo/client modules on API usage.

! instantiated-by: Prototype 3 — Graph-Memory API & Seed Graph [💡/络 🎯/已]

+context: 'apps/graph-memory' defines the Datascript schema, salience metadata, and store management functions that underpin FUTON1's fact graph. It coordinates in-memory state with XTDB persistence and is the central place where entity and relation documents are shaped before being written or read.

+if: You want FUTON1 to expose a clear, documented contract for how entities, relations, and salience fields are represented and mirrored between Datascript and XTDB, so that other futons and external tools can rely on the same facts without depending on internal implementation details.

+however: The existing README and tests describe schema and storage behaviour but do not yet present a concise API-level view of graph operations. Helper functions are still oriented toward REPL use, naming conventions vary, and there is no single description of what is considered part of the persistent "seed graph" or which documents are guaranteed to be mirrored into XTDB.

+then: Promote 'graph-memory' to a documented API: define the core entity and relation operations (init, upsert, link, query), state explicitly which fields and document types are mirrored into XTDB, and describe any default seed graph that is expected to exist in a fresh store. Align naming and function boundaries with this contract and treat incompatible changes as migrations rather than ad-hoc edits.

+because: A clear graph-memory contract turns FUTON1's internal schema and store logic into a reliable interface. This allows higher futons and external tools to depend on the same graph representation, reduces ambiguity about what is persisted, and makes changes to the underlying storage behaviour visible and deliberate.

+evidence:

each.

+because: Once goldens exist, small tune-ups stop breaking entire colony behaviour.

+evidence:

- 'test/resources/goldens/hunger_tau.edn' (163 lines): 6 golden sequences as specified.

- 'normal' : h=0.42→0.427→0.233, tau=2.076 (strong feeding effect).

- 'starvation' : h=0.88→0.896→0.887, tau FLOORS at 0.605 (exploitation mode).

- 'overstimulation' : h=0.50→0.506→0.262, tau PEAKS at 2.011 (exploration mode).

- 'heavy-cargo' : cargo=0.95, h=0.58→0.961, tau FLOORS at 0.597 (load stress).

- 'high-risk' : h=0.60→0.671→0.649, tau=1.14 (moderate modulation).

- 'high-ingest' : h=0.70→0.67→0.185, tau PEAKS at 2.18 (strong relief).

- 'test/ants/ai/hunger_regression_test.clj' (29 lines): Loads 6 scenarios, evaluates tick-hunger → update-hunger → hunger->tau → modulate-precisions, compares to golden.

- 'tools/hunger_goldens.clj' (75 lines): Generator script for golden file.

- 'README.md' (lines 117-129): Documents hunger config parameters and defaults.

+next-evidence:

- 'doc/hunger-goldens-guide.md' explaining behavioral regime each scenario represents.

- SVG plots for hunger_tau trajectories (like trace_tick.clj does for microtraces).

- Tau-floor/tau-cap assertions added to regression test (currently equality-only).

- Boundary saturation tests (e.g., h=0.99 + 0.1 delta → clamps to 1.0).

! instantiated-by: Prototype 4 — Policy Layer Surfacing [💡/乡 🐾/内]

+context: 'ants.ai/policy' selects actions using EFE metrics while 'attach-policy-diagnostics' emits traces.

+if: You want reproducible policy choices under controlled sensory inputs.

+however: There is no standalone harness that feeds synthetic mu/observation pairs.

+then: Build 'eval-policy mu prec obs config' as a harness and add golden action rankings for ten canonical scenarios (e.g. rich food, enemy close, starvation risk).

+because: Stabilised policies are mandatory before benchmarking or bridging into futon1 salience services.

+evidence:

- 'test/resources/goldens/policy_harness.edn' : 5 golden scenarios with action rankings.

- 'forage-normal' : Empty-handed, good food signal → forage preferred.

- 'loaded-homebound' : High cargo, near home → return preferred.

- 'white-space-scout' : No signals, exploring → pheromone preferred.

- 'deficit-colony' : Low reserves, high hunger → return prioritized.

- 'maintain-pheromone' : Good trail, near home → pheromone preferred.

Each includes: p (probability), G (free energy), risk, ambiguity, info, colony, survival, action-cost.

- 'tools/policy_harness.clj' (128 lines): Generator for 5 canonical scenarios.

- 'test/ants/ai/policy_harness_test.clj' (20 lines): Loads scenarios, runs 'eval-policy', compares rankings to golden.

- 'test/ants/ai/policy_test.clj' (187 lines): Comprehensive behavioral tests

(shape, preference coupling, starvation tau clamping, nest action filtering).

- 'src/ants/ai/policy.clj' : 'eval-policy()' function exists (deterministic ranking harness).

+evidence:

- Expand to 10 canonical scenarios as specified in devmap (currently 5).

- Document EFE component formulas (risk, ambiguity, info, colony, survival) in 'doc/policy-spec.md' .

+next-evidence:

- 'ants.ai/policy' now has a 'check!' API (inputs: pattern ID, context EDN, evidence refs; outputs: status, missing fields, derived tasks), enforce schema validation, and log every check as a proof state record.

+because: Turning patterns into executable checks is what makes FUTON3 "flexiformal" rather than just narrative.

+evidence:

- Checks today are conversational heuristics in aob-chatgpt or handwritten flexiargs.

- Define a 'check!' API (inputs: pattern ID, context EDN, evidence refs; outputs: status, missing fields, derived tasks), enforce schema validation, and log every check as a proof state record.

+because: This yields genuine pattern-aware intelligence.

! instantiated-by: Prototype 3 — Applicability Engine & Check DSL [💡/未 📈/已]

+context: We already use a "fake embedding" (sigil adjacency) to find nearby patterns.

+if: You want reproducible neighbourhood queries for pattern suggestions.

+however: The embedding lives in ad-hoc CSVs and isn't wired into FUTON3's APIs.

+then: Integrate the sigil-distance matrix into FUTON3 as a first-class service (pattern search API, 'nearest-patterns' command, CLI/UI views) and publish tests proving deterministic neighbourhoods.

+because: Pattern suggestion quality underpins the flexiformal workflow and must be inspectable.

! instantiated-by: Prototype 3 — Pattern Library Activation [💡/已 📈/世]

+context: You have dozens of patterns (Mojo, ORP, Paramita, Flexiarg, etc.).

+if: You want agents to recommend patterns or detect pattern-relevant situations.

+however: Patterns are not yet encoded in memory or linked to triggers.

+then: Select 2-3 patterns from the core set, develop them fully in one domain (e.g., ants or MetaCA), then apply them to a different domain (e.g., musical improvisation or crowdfunding) and document whether they work, fail, or require adaptation. Record the results as evidence for or against the transfer thesis.

+because: A key milestone for FUTON3's plausibility is demonstrating that patterns developed in one domain actually transfer to another—not just that they can be written, but that they work.

! instantiated-by: Prototype 9 — StackExchange Import [💡/及 📈/久]

+context: StackExchange holds thousands of Q/A pairs that fit Araxa's schema.

+if: You want to harvest them as production rules with provenance.

+however: there is no importer aligning their Markdown into our graph schema.

+evidence:

- 'src/futon3/checks.clj'

- 'src/f2/router.clj'

- 'src/f2/transport.clj'

(Futon1 bridge wiring), 'src/futon3/futon1_bridge.clj', and

- 'test/transport-test.clj' (workday/check forwarding cases) exercise the DSL end-to-end.

! instantiated-by: Prototype 5 — World Mechanics Contract (war.clj) [💡/书 📈/已]

+context: The simulator handles evaporation, movement, gather/deposit, pheromone dynamics, hunger propagation, reserves, and termination.

+if: You want FUTON2 to be a reliable agent-environment backend.

+however: Movement heuristics, gather thresholds, and white-streak rules are powerful yet lightly documented.

+evidence:

- Document all invariants (pheromone caps, grid bounds, gather limits), seed movement

documented or exercised a live-in-server workflow the way Emacs users rely on emacsclient.

+because: This becomes the living mathematical-philosophical substrate.

+next (inclusion/transclusion UX):

- Visually differentiate included vs. transcluded regions in the main article buffer.

+however: The REPL is off by default, and there is no canonical checklist proving that a Drawbridge session can connect, query MUSN state, and execute a minimal set of safe operations.

+then: Validate the Drawbridge workflow against a running Futon3 server: enable Drawbridge, connect via a REPL client, and execute a small script that queries 'musn.api/*' functions (clients/history/links/now), verifies read-only access, and records a short transcript in the logs/docs.

+because: The REPL is off by default, and there is no canonical checklist proving that a Drawbridge session can connect, query MUSN state, and execute a minimal set of safe operations.

@audience futon-devs, reflective-agents

@tone analytic-concise

@style flexiarg

+context: Hyperreal requires a consulting practice that increases alignment between your research, your livelihood, and your long-term direction while generating reusable structure rather than bespoke one-offs.

+if: You want each engagement to contribute both to financial stability and to the evolution of a sustainable, research-informed independent practice.

+however: Without intentional structure, communication slips; deliverables fragment into bespoke artefacts; and the insights generated fail to loop back into the system that needs them.

+then: Keep client communication transparent; produce deliverables as reusable, documented modules; and extract consulting insights as proto-patterns that flow into the Hyperreal backlog and FUTON3's salience pipeline.

+because: This turns consulting into an alignment engine: work generates reusable structure, structure generates higher variety, and higher variety increases your future alignment opportunities. Hyperreal becomes evolutionary rather than transactional.

+evidence:

- Existing consulting deliverables already include modular components.

- Client communication typically happens through channels that are easily timestamped and reviewable.

- Notes from past engagements show extractable conceptual structure.

- The VSATLAS work demonstrates a clean example: a reusable flexiarg design that extends the client's system while revealing future service-provision revenue streams.

+next: Introduce a lightweight "module template" for new deliverables; add a brief reflection step at the end of each engagement; and feed extracted insights into FUTON3 for proto-pattern surfacing.

+next-evidence:

- Detection of repeated module reuse across engagements.

- FUTON3-generated proto-patterns based on recurring consulting moves.

- Evidence of alignment gain: smoother scoping, clearer pricing, better-fit clients, and reduced bespoke overhead.

**! instantiated-by: Prototype 5 — Ambition 2 FUTON-Araxana Stack Stewardship [❤️/用
@@/久]**

+context: FUTON0 needs a reliable, steadily improving environment for thought. Daily use and clear architectural notes keep the FUTON-Araxana stack coherent and usable.

+if: You want FUTON-Araxana to become a dependable workspace for reasoning, development, and collaboration.

+however: Skipped daily use weakens the system; undocumented decisions cause drift; and premature collaboration introduces scope creep.

+then: Use FUTON1 daily as the live testbed; record architecture reflections as they arise; and add collaborators gradually with guardrails that protect scope and focus.

+because: Frequent contact and lightweight documentation prevent drift and make the system stronger over time. This keeps FUTON-Araxana evolving into a stable, coherent environment rather than an abandoned prototype.

+evidence:

- FUTON1 already supports daily usage with minimal friction.

- Architecture notes demonstrate where documentation has prevented drift.

- Early collaborative experiments (e.g., with Rob) show the value of guardrails.

+next: Stabilise daily usage with light scaffolding; formalise a minimal architecture note-taking protocol; and extend Araxana gradually as stable components emerge.

+next-evidence:

- Daily FUTON1 usage detectable via timestamps.

- Architecture notes showing reduced drift or clearer decision points.

- Araxana prototypes yielding stable nodes or literate fragments.

**! instantiated-by: Prototype 6 — Ambition 3 Mojo, Evolving Variety, + Reflective Rhythm [❤️/弓
👉/又]**

- README documents the Datascript schema, XTDB touchpoints, and hydration behaviour used by FUTON1.

- Test suites for graph-memory exercise persistence and salience fields.

- XT inspection paths demonstrate how entity and relation documents are stored and reloaded.

+next: Extract and name the core graph-memory operations as a public API, document which document types and fields are mirrored into XTDB, and describe any expected seed graph in the module README.

+next-evidence:

- An updated 'apps/graph-memory' README that presents the API, mirroring rules, and seed-graph expectations in one place.

- Tests that target the documented API functions rather than only internal helpers.

- A short migration note for any changes that affect stored documents or XTDB layout.

**! instantiated-by: Prototype 4 — XTDB Operations & REPL Workflows [🔍/分
👤/今]**

+context: FUTON1 mirrors all graph mutations into XTDB and provides a stable hydration path so that entity and relation documents, salience fields, and type data persist across restarts. The repository already includes REPL notes, inspection snippets, and a dedicated '-M:storage/inspect' entry point for examining the store directly.

+if: You want developers to inspect, query, and verify the persistent store using a small set of predictable entry points, without relying on ad-hoc XTDB queries or bespoke REPL sessions.

+however: Common inspection tasks—listing entities, examining relations, checking salience fields, or understanding schema layout—still require manual XTDB queries or scanning several snippets across READMEs and test files. These workflows are reliable but not yet consolidated into a minimal operational interface.

+then: Package the existing XTDB inspection workflows into a small set of reusable functions (e.g., list-entities, list-relations, inspect-entity, counts), present them in one place with brief examples, and highlight the salience-survives-restart demonstration as a canonical illustration of the hydration contract. Treat these recipes as the reference interface for inspecting and understanding FUTON1's persistent state.

+because: Consolidating XTDB operations into clear, reusable workflows reduces friction, makes store inspection more consistent, and strengthens FUTON1's role as a deterministic storage substrate rather than an opaque backend. Developers gain a predictable way to understand how mutations appear in XTDB and how hydration restores them.

+evidence:

- 'README-storage.md' documents XTDB inspection behaviour and usage.

- The '-M:storage/inspect' alias provides automated examples of querying the store.

- Hydration and salience persistence are demonstrated through existing scripts and XT touchpoints.

+next: Surface the same XTDB schema and inspection workflows in Emacs so FUTON3 can browse persistent store contents without relying on filesystem paths or ad-hoc REPL interaction.

+next-evidence:

- A simple Emacs view that displays XT entities, relations, and salience metadata using the documented inspection functions.

- Consistent behaviour between the Emacs view and the REPL inspection entry points.

**! instantiated-by: Prototype 5 — Open-World Ingest Workflows [📝/开
👤/人]**

+context: 'apps/open-world-ingest' streams arbitrary text through a deterministic CoreNLP pipeline and writes entities, mentions, and relations into XTDB. These documents become part of the same persistent store that the demo/client workflow hydrates on startup, allowing background knowledge to accumulate outside interactive sessions.

+if: You want FUTON1 to support both offline corpus ingestion and interactive exploration, with both surfaces reading from and writing to the same XTDB state in a predictable way.

decisions for determinism, and split movement modes into testable units.

+because: A clear contract keeps the world trustworthy when other prototypes depend on it for episodes.

**! instantiated-by: Prototype 6 — Live Analyzer & Pivot Stream Stabilisation [💻/E
🌐/无]**

+context: The live analyzer consumes pivot events, computes EWMA metrics, detects starvation spirals, and prints mode transitions.

+if: You want analyzer outputs to feed futon3 evaluations or futon4 reflexive agents.

+however: Pivot schemas are implicit and analyzer assumptions are undocumented.

+then: Freeze a versioned pivot-row schema, add tests for required keys (:act, :mode, :cargo, :ing, etc.), and ship an example analyzer plugin for futon4.

+because: Pivot streams are FUTON2's event bus and must behave like a stable ABI.

**! instantiated-by: Prototype 7 — Classic vs AIF Benchmark Suite [💻/三
👤/不要]**

+context: Benchmark scripts compare classic ants to AIF runs across scores, durations, and G-ema.

+if: You want formal behavioural comparisons rather than anecdotes.

+however: Only three preset scenarios exist and there is no variance analysis.

+then: Add parametrised benchmarks spanning grid sizes, pheromone decay, hunger burn, and tau floors; capture plots and golden CSVs.

+because: FUTON2 should function as a scientific system with reproducible metrics.

**! instantiated-by: Prototype 8 — External Control & Episode Export [🤖/予
💻/已]**

+context: 'simulate' drives the war sim but offers no external stepping API.

+if: You want futon3 or futon4 to request "give me N steps with this config."

+however: Stepping remains tied to the CLI and HUD.

+then: Expose '(futon2.step world n {emit-pivot? true})', add EDN exports of full episodes, and keep pivot batches consistent.

+because: Exported episodes are critical for agent training, memory, and pattern extraction.

**! instantiated-by: Prototype 9 — futon1 ↔ futon2 Alignment Layer [💡/冂
👤/二]**

+context: FUTON1 emits focus headers and salience; FUTON2 emits percept/action streams.

+if: You want the ants to act as embodied readers + writers of futon1's graph.

+however: There is no shared dictionary between entities/intents and colony modes, so exports are useless to FUTON1.

+then: Define the crosswalk (entities ↔ colonies, intents ↔ behaviours), publish an export schema, and prove it by replaying one colony while FUTON1 ingests the stream and surfaces it back to users.

+because: a working bridge proves ants can reason over the same facts as the rest of the stack.

**! instantiated-by: Prototype 10 — Curriculum Scenarios & AIF Behaviour Cards [📖/辻
👤/予]**

+context: The current war scenario is rich but single-purpose.

+if: You want FUTON2 to double as a pedagogical engine for Active Inference.

+however: There is no named curriculum.

+then: Create behaviour cards (home-return, starvation spiral, over-exploration, defensive lattice, greedy gatherer collapse), simulate each with fixed seeds, and export the episodes.

+because: These cards seed pattern extraction and future training systems.

**! instantiated-by: Prototype 11 — AIF Agent Debugger & Mu/Precision Visualiser [🤖/由
💻/-]**

+context: Pivot logs are textual and require manual parsing.

+if: You want interactive introspection of mu vectors, tau evolution, and weighted errors.

+however: No visualiser or "follow Alice" mode exists.

+then: Build a Clerk notebook (or similar) that plots tau, hunger, errors, and pivot timelines, plus a mode that traces one agent's full mu timeline.

+because: Visual insight accelerates iteration and debugging.

! instantiated-by: Prototype 12 — Year-End Packaging & Story [📝/未]

'obligation/id', ':action/tags', and ':delta/joy', add rollups that show which devmap clauses advanced per day, and expose a '/musp/trails/proof' export for futon4 archives.

+because: The whole point is to measure how daily practice advances the devmaps.

+design: The Tatami HUD cue-reflection pattern mirrors every FROM-CHATGPT entry back through '/musp/cues' so fruits/orbs in the HUD match the transport's cue embedding ('comtrib/aob-chatgpt.el:1503', 'src/f2/ui.clj:97', 'src/futon3/cue_embedding.clj:214').

+evidence: 'resources/tatami-events.edn' captures live tatami sessions, 'resources/tatami-context.edn' logs every FROM-CHATGPT-EDN HUD payload, and the cue embedding tests ('test/futon3/cue_embedding_test.clj') prove keyword/intentionally-null payloads round-trip.

+blocker: [/trail/persistence:cue-validation]

+next: (1) persist cue annotations into the trial export so Prototype4 records which fruits/orbs were actually computed, and (2) add proof artifacts (HUD screenshots or '/musp/cues' transcripts) that demonstrate the cues signal real changes instead of static defaults before declaring the prototype stable.

**! instantiated-by: Prototype 12 — Year-End Synthesis [📝/未
🌐/轮]**

+context: The prototype cycle ends; reflection and forward-planning needed.

+if: You want long-term stability, and upgraded beta-schema.

+however: Accumulated divergence between prototypes must be unified.

+then: Run full PAR + CLA + DREAM integration cycle.

+because: This sets the stage for FUTON5+.

improvisational rather than dogmatic.

+however: Without review, patterns accrete into heaviness.

+then: Run a yearly cycle that prunes dead patterns, refines musical grammars, updates CT riffs, consolidates improvisation patterns, refreshes the Nonstarter ethos, and removes anything that smells like KPIs.

+because: Concentration equals aliveness, and aliveness must be renewed.

+context: FUTON0 needs a source of evolving, meaningful variety. Moments of “wondrous apprehension”—small encounters with insight, novelty, or strangeness—keep the internal model flexible and alive.

+if: You want learning, creativity, and self-understanding to deepen through controlled contact with the unknown rather than through routine or ritual.

+however: When these encounters vanish or are postponed indefinitely, patterns ossify, curiosity shrinks, and alignment stops increasing.

+then: Invite small, natural opportunities for wondrous apprehension—through ideas, conversations, creative acts, or symbolic cues—and metabolise them lightly, without turning them into obligations or scheduled rituals.

+because: These brief encounters act as evolutionary impulses: they expand perspective, refresh the generative model, and increase adaptive variety. This introduces just enough uncertainty to stimulate growth without destabilising the system.

+evidence:

- Rhythm experiments such as PAR sessions, RAP sessions, Mojo prompts, and card draws have reliably triggered moments of insight without becoming rigid.

- Creative or relational encounters have historically produced generative perspective shifts.

- Minimal reflection has been sufficient to integrate these signals.

+next: Keep the system open to small shocks of insight; capture only what feels naturally significant; and allow FUTON1 or FUTON3 to mark these moments lightly when they matter.

+next-evidence:

- Occasional tags noting spontaneous insight or perspective shift.

- Detectable reductions in stagnation or pattern ossification.

- FUTON3 identification of proto-patterns emerging from these encounters.

! instantiated-by: Prototype 7 — Ambition 4 Role Transition Trajectory [↗/口/心/用]

@multiaig futon0/prototype7
@title Ambition 4 — Applied Futon Theory: Co-Evolution With Future Environments
@audience futon-devs, reflective-agents
@tone analytic-concise

@style flexiarg

+context: FUTON0 must support your co-evolution with future environments—roles, collaborators, and communities that do not yet fully exist but already exert directional influence as futons (real-but-virtual futures).

+if: You want your future working life to emerge through mutual shaping rather than reactive job-seeking or inherited scripts.

+however: When futures are treated as fixed targets or when symbols are mistaken for the futons they mediate, misalignment grows and trajectories become brittle.

+then: Attend to symbolic signals that reveal possible future environments; explore lightly; cultivate relationships that show reciprocal responsiveness; and let alignment emerge through ongoing co-evolution rather than forced planning.

+because: Futons influence the present something like photons. The co-evolution of present and future echoes Gotama’s dependent causality—agent and future conditions arise together in a logical chain—and Alan Moore’s insight that all times are immediately accessible in certain states of consciousness. In applied futon theory, alignment grows through co-evolution with virtual states, not through selection or prediction.

+evidence:

- Prior trajectories (Hyperreal, VSATLAS, ORP, band work) emerged through reciprocal adaptation, not explicit search.

- Symbolic cues—patterns, conversations, prototypes—have repeatedly revealed access to futures not yet actual.

- Misaligned contexts typically show low responsiveness early, acting as non-resonant futons.

+next: Track signals of futonic resonance or dissonance; maintain light exploratory contact with promising environments; and let the trajectory refine itself through reciprocal influence rather than pressure.

+next-evidence:

- FUTON1 tags capturing futonic moments

+however: Ingest currently operates as a standalone CLI with limited guidance on how its outputs integrate with the demo/client view of the world. Directory conventions, expected data-root layouts, and test coverage for ingest semantics (ID stability, relation structure, negation handling) are not yet consolidated.

+then: Define a standard data-root layout for ingest and demo/client workflows, document how ingest-produced XTB documents appear in the hydrated store, and provide small example workflows that ingest a text sample and then query the resulting state through the demo/client or API surfaces. Add a minimal set of golden tests covering entity-ID stability and relation formation.

+because: When ingest and interactive exploration share a clear, documented storage path, FUTON1 can accumulate durable background knowledge in a controlled way. This makes the persistent world state coherent across offline and online workflows and maintains determinism across restarts.

+evidence:

- README confirms deterministic ingest via CoreNLP and XTB mirroring.

- Demo/client workflows hydrate from XTB on startup and therefore can read ingest-produced documents.

- Existing golden fixtures demonstrate deterministic behaviour in other pipeline components.

+next: Document shared data-root conventions; provide an ingest+query example in the README; and implement golden tests for ingest ID stability and relation structure.

+next-evidence:

- A worked example showing ingest followed by demo/client queries against the same XTB store.

- Golden tests verifying stable entity IDs and relation outputs for a fixed input.

- README updates linking ingest workflows to the persistent world model.

! pending: Prototype 4 — NLP Experiments & Traceable Knowledge Graphs [起/現/知]

+context: The new CLI tracer ('scripts/demo_trace.cl') and API tap listener ('api.analytics/start!') produce structured focus-header and profile payloads for every turn. These traces make NLP regressions visible and allow us to reconstruct session-specific ER graphs outside XTB.

+if: You want to compare alternative recognisers or mathematical-text pipelines against the deterministic substrate, or you want to convert Tatami/demo sessions into inspectable knowledge graphs without mutating production storage.

+however: Tatami tracing is still manual, and there is no tooling yet to turn the tap payloads into graph artifacts. Ideas for richer NLP experiments are not yet captured in this devmap.

+then: Extend the headless API/tatami flow with the same tap logging as the demo tracer, document the planned NLP experiments here, and build a small exporter that reads trace EDNs + tap events and emits a consolidated entity/relation snapshot per session.

+because: Consistent traces across CLI and API let FUTON1 evaluate NLP changes deterministically and produce reusable “knowledge graph” snapshots per session. Recording these plans keeps the work visible without derailing core storage duties.

+next:

- Add an opt-in tap sink to the API/Tatami pipeline so live sessions populate '~code/storage/futon1-traces/<session>' ‘just like the demo tracer.’

- Write an ER graph exporter that ingests those traces and produces a summary artifact (EDN/Graphviz) for analysis.

- Catalogue candidate NLP experiments (alternate recognisers, mathematical text ingestion, scoring heuristics) and their evaluation criteria in this devmap.

+next-evidence:

- Trace directories containing both CLI and Tatami sessions.

- ER graph export script plus sample output checked into ‘resources’ or ‘docs’.

- devmap entries describing the queued NLP experiments and how their results will be judged.

+instantiated-by: Prototype 6 — Thematic Clusters & Intent Metadata [💡/无/助]

+context: FUTON2 must read as a coherent simulation layer to partners and future prototypes.

+if: You want it to be adoptable as a research simulator.

+however: The codebase is powerful but intimidating without a unifying story.

+then: Package the release with a top-level README, whitepaper, stepping API, goldens, and versioned pivot schema.

+because: Strong packaging turns FUTON2 into a reusable, inspectable module inside the full FUTON stack.

! instantiated-by: Prototype 13 — Graph Memory Adapter & Agent ABI [💡/调]

+context: FUTON1 now offers a documented graph-memory API; FUTON2 needs to speak it to graduate beyond the sandbox.

+if: You want ants to be the first drop-in “AI agent” that reads fact graphs and writes back its beliefs.

+however: perceptions/actions remain locked to grid coordinates, so nothing maps to FUTON1 entity IDs or relation semantics.

+then: Design a bidirectional adapter (percepts → graph facts, graph cues → priors), version the API, and replay a colony run that exports beliefs into FUTON1 salience tables. Treat FUTON3 proof trails as “human pheromone lines”: export the canonical trail EDN (pattern IDs, fruits/orbs, joy deltas) into FUTON1, expose them via the adapter so analyzers ingest them as missions, and log agent runs as trail-style payloads (mission, outcome, energy delta) back into MUSN so FUTON3 records “agent proofs.”

+because: Once the adapter honours shared trail/pheromone energy, human proofwork and AI runs guide each other through the same mission record instead of bespoke glue.

! instantiated-by: Prototype 14 — Viriya Field & Social Activation Metrics [⚡/现/动]

+context: FUTON2 carries the Factor of Awakening “energy/virya”: vigor, commitment, and kinetic practice.

+if: You want simulations to reflect and inform real social activation—networked collaboration, shared commitments, and practice loops across humans + agents.

+however: Current metrics stop at hunger/tau; nothing measures vigor, pledge strength, or coordination throughput.

+then: Introduce viriya indicators (e.g., colony activation score, pledge adherence, transfer of trials into collaborative tasks), log them per episode, and export them so futon3/futon7 can read “energy health.”

+because: Tying the simulation to viriya turns FUTON2 into a practice accelerator instead of a curiosity.

! instantiated-by: Prototype 15 — Sandbox Bridge & Presence Interface [🌐/内/刊]

+context: FUTON3 now emphasises flexiformal proofwork but still hosts multi-user rooms and SAFE REPL affordances.

+if: You want AI agents (ants and successors) to inhabit those rooms, receive room/talk updates, and emit actions back into the simulation.

+however: No interface maps MUSN room events/presence into FUTON2 percepts, nor do ants announce themselves inside FUTON3.

+then: Define a bidirectional “presence bus” (rooms → observation cues, agent actions → MUSN events), add authentication so only designated colonies bridge, and ship a demo where ants report discoveries into a FUTON3 trail.

+because: This proves that AI agents can show up wherever proofwork happens, not just inside a war grid.

! instantiated-by: Prototype 16 — Lobby Bus & MUSN Intake [💡/无/已]

+context: FUTON3 runs with MUSN as a shared graph and chatgpt-shell as a HUD, but ChatGPT-API currently has only conversational access and no principled way to commit events back into MUSN.

+if: You want ChatGPT-API to propose or make “persistent” updates (breach logs, paramita orbs, devmap notes, pattern instantiations) without giving it full god-mode over the stack.

+however: There is no explicit one-room “Lobby” where HUD traffic is normalised into MUSN events, no single-writer policy, and no clear schema for what counts as an escalated, storables event versus transient chatter.

+then: Create a dedicated Lobby room in MUSN

“indicator-construction steps that should land in the blue virtue band”.

+next: (1) Generalise the existing cue embedding to operate on ‘trail/intent’ + ‘pattern/reasoning’ tuples and return (fruit, orb) as proof step types, (2) run a calibration pass over a small set of manually typed patterns to tune the fruit/orb basis and salience weights, ensuring that canonical relations like 🚶→🌐 (“value-aligned indicator checks”) land in the intended region, and (3) wire the enriched trail export into Prototype 4’s rollups so devmap progress can be inspected as “which constructive proof-types actually fired today, how salient they were, and whether the checks that claimed to align with your values really did so in the embedding.”

! instantiated-by: Prototype 5 — Workday Instrumentation & aob-chatgpt Bridge [🌐/每]

+context: Within ‘aob-chatgpt’ we already experiment with reflecting on daily actions against devmap obligations.

+if: You want FUTON3 to serve as the backend for those reflections.

+however: There is no official API for submitting “workday claims,” mapping them to patterns, or receiving verification feedback; roles of “workday participants” and “instrumentation stewards” are undefined.

+then: Publish a ‘workday/submit’ endpoint (inputs: timestamp, activity, evidence links), run it through the check DSL + embedding, respond with pattern hits/misses plus follow-up obligations, and name the actors explicitly (participants supply evidence; instrumentation stewards curate dashboards); integrate the same flow into the ChatGPT tooling.

+because: This is how empowered action (“joy”) becomes measurable proof progress and how role responsibilities stay legible.

+next: Replace the temporary ‘futon3/logs/workday.edn’ queue with a FUTON1 persistence call (‘workday->graph’) and add HTTP affordances + dashboards so instrumentation stewards can audit submissions without tailing logs.

+evidence: Workday intake lives in ‘src/futon3/workday.clj’ with coverage in ‘test/transport_test.clj’ (workday submission) and README sections describing the API.

! instantiated-by: Prototype 6 — Pattern Creation Workbench [💡/已]

+context: New patterns emerge from trails and devmap edits, but the capture is currently manual in futon4.

+if: You want FUTON3 to draft candidate patterns/flexargs from accumulated proof states.

+however: Manual backlog in FUTON4 creates tension; there is no readiness window for when drafts are “hand-off ready.”

+then: Build a summariser that clusters trails, extracts hooks, and emits draft flexargs blocks (IF/HOWEVER/THEN/BECAUSE) tagged with suggested paramita, ready for futon4 editors, and declare readiness criteria (e.g., at least N supporting checks + steward review) so handoffs are predictable.

+because: FUTON3 should not just consume the library; it should propose expansions grounded in lived work while signalling when drafts can be trusted.

! instantiated-by: Prototype 7 — Joy-as-Empowered-Action Metrics [💡/力/已]

+context: Joy here = Spinoza increase in power to act, evidenced by commitments you can now keep.

+if: You want metrics that show whether checks are empowering or draining.

+however: Current “joy metrics” counted UI friction; they did not measure proofs discharged or obligations unblocked.

+then: Track ‘time-to-proof’, ‘obligations-cleared’, ‘follow-on commitments’, and ‘virya linkages’ per session, and visualise them alongside subjective notes.

+because: Empowered action is the KPI for this futon.

! instantiated-by: Prototype 8 — futon2/futon1 Proof Hooks [💡/二/已]

+context: FUTON3 already generates proof states; FUTON1 (graph) and FUTON2 (energy) need those proofs as fuel.

+if: You want every accepted check to become a graph relation and a virya delta without manual glue.

+then: Create a dedicated Lobby room in MUSN

(resonance, responsiveness, momentum).
- Notes marking when symbolic cues open or close access to future attractors.
- Increasing clarity, lightness, and alignment as co-evolution proceeds.

! instantiated-by: Prototype 8 — Stack Coverage Check [💡/📝/?] [!] [!]

+context: Devmaps belong to FUTON3: they are symbolic membranes through which FUTON3 perceives futons and maintains structural coherence across the stack. FUTON0's role is to support the mindful conditions that allow these symbolic structures to appear, evolve, and be refreshed.
+if: You want FUTON3 to notice where the system's symbolic interfaces to future environments are missing, stale, or drifting, before the whole system loses alignment.
+however: When narrative, reflection, and wondrous apprehension are weakened, futons become harder to perceive; devmaps stagnate; and FUTON3 loses access to future-relevant structure.
+then: Maintain a light narrative awareness of the system's evolving futures; allow futons to be apprehended through conversation, reflection, and symbolic cues; and let FUTON3 formalise these perceptions into updated devmaps when needed.
+because: FUTON0 provides the experiential ground—a mindful openness to futons—while FUTON3 handles representational coherence. Without the former, FUTON3 cannot see; without the latter, FUTON0's perceptions cannot stabilise. Together, they keep the stack aligned across time.
+evidence:

- Narrative reflection in this conversation has repeatedly revealed futons that were not yet encoded in devmaps.
- FUTON3 updates gain clarity when experiential insight precedes them.
- Wondrous apprehension consistently precedes formalisation.

+next: Support FUTON3 by staying attuned to emerging futons; mark futon moments lightly when they arise; and trust FUTON3 to handle structural consolidation.

+next-evidence:

- FUTON1 entries capturing futon moments.
- FUTON3 updates that align with experiential insight.
- Clearer, more resonant devmaps appearing naturally over time.

! instantiated-by: Prototype 7 — CLI Ergonomics & Multi-Instance Isolation [🌐/资料/互/] [!] [!]

+context: The demo/client workflow exposes bang commands, flags, and environment variables for selecting data directories, overriding XTDB configuration, and controlling persistence behaviour. These surfaces already allow multiple instances of FUTON1 to run side-by-side on the same machine.
+if: You want developers and tools to run isolated demo/client or API sessions for testing, experimentation, or production-like scenarios without interfering with each other's storage or configuration.
+however: Best practices for using 'BASIC_CHAT_DATA_DIR', XTDB resource overrides, 'config.edn', and reset/compact/export options are spread across READMEs and command-line examples. It is easy to misconfigure data roots or XTDB paths, and no single reference describes the recommended conventions for running multiple instances safely.
+then: Document a small set of configuration patterns (e.g., dev/test/prod data roots), clarify how env vars and config.edn interact, and provide brief examples of running multiple isolated instances. Add smoke tests that launch separate demo/client or API sessions with distinct data roots. Include a simple '/status' helper or equivalent CLI flag that reports the active configuration.
+because: Clear ergonomics and isolation practices make FUTON1 safer to embed in tooling and CI workflows, reduce accidental data mixing, and strengthen the determinism and reproducibility guarantees laid out in FUTON1's IFR.
+evidence:

- README documents environment overrides, config.edn usage, and the structure of XTDB data directories.

+context: FUTON1 can attach lightweight thematic tags to utterances or text spans during ingest or interactive use. These tags—derived from deterministic NLP features or simple pattern-recognition hooks—provide a minimal layer of "intent metadata" that downstream futons may use for analysis or policy.

+if: You want FUTON1 to expose a small, deterministic set of thematic categories that describe common conversational moves or content types, without relying on probabilistic classifiers or embedding higher-level inference in this layer.

+however: There is no consolidated thematic taxonomy, the tags used in logs are not documented as part of the public contract, and no examples show how these tags appear in the stored XTDB documents or hydrated in-memory state. Integration points within demo/client or ingest workflows remain implicit.

+then: Define a minimal set of thematic clusters supported by deterministic NLP features, document them as part of FUTON1's public contract, and ensure that produced tags (when enabled) appear predictably in both in-memory and XTDB representations. Provide a small set of example workflows illustrating how these tags are attached during ingest or interaction and how they can be queried or inspected.

+because: Lightweight thematic metadata strengthens FUTON1's role as an inspectable storage substrate. By keeping the taxonomy small and deterministic, FUTON1 avoids embedding inference or policy while still providing structured signals that higher futons may use for reasoning or pattern-based behaviour.

+example: To assist the user with time management, ChatGPT-API offered: "Do you want breaches auto-logged if you're still active at 22:10?" At present there's no way for ChatGPT to "actually" log such breaches anywhere. Rather than assuming all escalated events must be elicited in the same way at all times, allow 'set-user-system-prompt' in futon3's 'aob-chatgpt.el' to switch between modes (e.g. ':normal', ':breach-logging') under external control (cron, or a MUSN agent). ChatGPT-API can then "propose" mode shifts or breach events via the Lobby, while the MUSN side decides which proposals to enact.

+example2: While reviewing a day's evidence, ChatGPT-API proposes a structured "paramita reflection" event via the Lobby:

```
{:musp/events
[{:event/type :paramita/reflection
:event/paramita :equanimity
:event/notes "User can apply 9 of Cups to understand their own mind and body."}
{:event/request :musp.store!}]}
```

This safely introduces reflective practice artefacts into MUSN without conflating them with actual mode changes or breach logging. It demonstrates that the Lobby is not only

for operational concerns (like curfew rules) but also for epistemic/ritual artefacts that support the broader FUTON practice space. The HUD then renders a subtle, clickable banner: Reflection exercise suggested >

! instantiated-by: Prototype 16 — Wisdom Trail Harvest & Pattern Seeding [⌚/去💡/氏] [!] [!]

+context: FUTON3 trails now record which patterns were checked or advanced during human work sessions.

+if: You want FUTON2 analyzers to ingest those trails, compare them to colony behaviour, and commission new agent goals.

+however: There is no ingestion path for trail EDN, nor a feedback loop that says "agents, go gather evidence for Pattern X."

+then: Add a 'trail->mission' translator that turns FUTON3 proof deltas into FUTON2 goal queues, plus a reporting job that shows how agent runs affected subsequent checks.

+because: Shared missions keep the "energy" loop tight between simulations and human proofwork.

+if: You want the system to describe itself in a way that agents can query, and for that description to stay synchronized with implementation.

+however: No machine-readable manifest exists where components are first-class entities in the hypergraph. Patterns describe design structures but not system capabilities.

+then: Unify documentation and APIs into a machine-readable manifest where stack components are first-class entities in FUTON1. Patterns describe not just design structures but system capabilities—what the message bus can do, what functions the Emacs interface exposes, what invariants each layer enforces. Enable agents to query the system about its own capabilities, and self-validate that the running system matches its declared structure.

+because: This closes the reflexivity loop: a semantic network that includes a model of itself. FUTON3's pattern checker already validates its own devmap; extending this to describe the full stack enables ongoing self-documentation as the system extends. The machinery from F5 (patterns → schema → code) applies reflexively: the stack itself becomes a domain.

+evidence:

- 'futon4/docs/evidence/rc-1.1-evidence.edn' seeds a machine-readable claim map for FUTON0-FUTON4.

+next:

- Expand the evidence map into a manifest schema (components, interfaces, invariants) and ingest it into FUTON1 for queryable self-description.

+however: today the exports are bespoke Clerk notebooks or one-off scripts.

+then: standardise two adapters—'proof->graph' (materialise relations with provenance) and 'proof->energy' (emit activation deltas for futon2 dashboards)—and lock them with fixtures + schemas.

+because: proofs only matter when other futons can act on them automatically.

+next: Wire the Tatami event log into FUTON1 (per 'futon3/tatami_store.clj' TODO) so session events flow into graph-memory before we expose the adapters.

! instantiated-by: Prototype 9 — Proof Training Ground ("Flexiformal Agent Sandbox") [🤖/汎/未] [!] [!]

+context: FUTON3 needs a reproducible training ground where humans and agents can practice flexiformal checks before contributing to FUTON6.

+if: You want collaborators (including FUTON6 math stewards) to experience the proof workflow hands-on.

+however: README/docs still emphasise ants and MUD metaphors and do not ship a simulator/tutorial bundle; roles (training facilitators vs. FUTON6 learners) and release readiness are not declared.

+then: Package a "Flexiformal Training Ground" release: updated README, tutorials referencing 't4r' training patterns, sample check sessions, and a minimal simulator that feeds FUTON6's Hyperreal entries; name facilitators (FUTON3 maintainers) and consumers (FUTON6 trainees), and publish readiness criteria (tutorials complete, simulators seeded, monitoring facets updated) before tagging the release.

+because: Storytelling plus an explicit training ground turns FUTON3 into a reproducible service that upstream layers (FUTON6, FUTON7) can consume and steward.

+pattern: Adopt 't4r/supplement' so the training ground demos mirror the Supplement workshop pattern—multi-agent listening infrastructure, governance rehearsal, and design-tension prototyping—anchoring the release narrative in the same staging choreography.

+next: Aggregate Tatami session proof events into 'proof-chain.edn' (per 'tatami_store.clj' TODO) so the training ground exports persistent transcripts.

! instantiated-by: Prototype 10 — System Self-Description & Machine-Readable Manifest [🌐/本/互/系统] [!] [!]

+context: The stack currently has documentation and APIs at various layers—FUTON1's storage API with checked invariants, Emacs Lisp functions callable via emacsclient, HTTP/WebSocket endpoints in FUTON3. But these exist as separate artifacts: code, docs, and schemas that humans must correlate.

+if: You want the system to describe itself in a way that agents can query, and for that description to stay synchronized with implementation.

+however: No machine-readable manifest exists where components are first-class entities in the hypergraph. Patterns describe design structures but not system capabilities.

+then: Unify documentation and APIs into a machine-readable manifest where stack components are first-class entities in FUTON1. Patterns describe not just design structures but system capabilities—what the message bus can do, what functions the Emacs interface exposes, what invariants each layer enforces. Enable agents to query the system about its own capabilities, and self-validate that the running system matches its declared structure.

+because: This closes the reflexivity loop: a semantic network that includes a model of itself. FUTON3's pattern checker already validates its own devmap; extending this to describe the full stack enables ongoing self-documentation as the system extends. The machinery from F5 (patterns → schema → code) applies reflexively: the stack itself becomes a domain.

+evidence:

- 'futon4/docs/evidence/rc-1.1-evidence.edn' seeds a machine-readable claim map for FUTON0-FUTON4.

+next:

- Expand the evidence map into a manifest schema (components, interfaces, invariants) and ingest it into FUTON1 for queryable self-description.

- Demo/client and API workflows can already be run with separate data roots.
- XTDDB instances behave independently when configured with distinct paths.
+next: Provide documented configuration patterns, add isolation-oriented smoke tests, and implement a '/status' helper that surfaces the active config.
+next-evidence:
- A README section showing dev/test/prod data-root conventions.
- Automated smoke tests confirming independent operation of concurrent instances.
- A '/status' output (CLI or API) showing current data-root and XTDDB config.

! instantiated-by: Prototype 8 — Golden Scripts & Determinism Checks [!

+context: FUTON1 already uses scripted EDN conversations and golden outputs to verify deterministic behaviour in the demo/client workflow, the NLP interface, and graph-memory operations. These fixtures ensure that focus-header formation, XTDDB hydration, and pipeline flow remain stable across revisions.

+if: You want confidence that changes to FUTON1 do not introduce regressions in determinism, header structure, or end-to-end behaviour, and that core workflows continue to perform within reasonable runtime bounds.

+however: Golden coverage remains incomplete. Existing fixtures focus on specific modules rather than end-to-end flows, and there is no consolidated suite capturing short conversational runs, longer histories, or ingest-driven scenarios. Runtime behaviour is observable but not yet checked in a structured way.

+then: Curate a small suite of canonical golden scripts that exercise representative workflows: a short demo/client session, a longer session with history, and a minimal ingest→query example. Record their expected outputs and basic runtime characteristics, and run them under CI so deviations in XTDDB state, focus-header structure, or runtime ranges are immediately visible.

+because: A compact, well-chosen golden suite makes FUTON1's determinism guarantees explicit, stabilises expectations about performance, and turns the substrate into a measurable layer rather than an informal demonstration.

+evidence:

- README documents deterministic behaviour, hydration patterns, and golden fixtures already present in module tests.
- NLP and graph-memory test runners use scripted inputs and expected outputs.
- XTDDB inspection tooling ('-M:storage/inspect') demonstrates reproducible state across runs.

+next: Assemble the canonical golden suite, add runtime checks under CI, and link the suite from the top-level README as FUTON1's determinism reference.

+next-evidence:
- A small directory containing canonical scripts and their expected outputs.
- CI jobs reporting on determinism and basic runtime stability.
- README references explaining how to run and interpret the golden suite.

! instantiated-by: Prototype 9 — External Client & API Contract Clarification [

+context: FUTON1 already exposes an HTTP API ('apps/api') that lets external tools query and update the persistent store through the same command-service logic used by the demo/client workflow. This provides a minimal surface for submitting utterances, inspecting state, and retrieving deterministic summaries such as focus headers.

+if: You want external clients to interact with FUTON1 without scraping CLI output or depending on internal modules, and to rely on a stable, documented interface for obtaining focus-header data and querying the underlying graph.

+however: The available endpoints and their expected inputs and outputs are documented across several READMEs rather than in a single reference, and there are no minimal example clients showing how to issue requests and interpret the deterministic responses. The boundary between FUTON1's API and higher-level agent behaviour remains implicit.

+then: Consolidate FUTON1's HTTP API contract:

--	--	--	--	--

document the endpoints, describe the structure of request and response bodies (including focus headers), and provide one or two minimal example clients that demonstrate how to submit an utterance and inspect the resulting state. Keep the API minimal and deterministic, leaving policy or agent behaviour to higher futons.

+because: A clear, documented API surface allows external tools to rely on FUTON1's deterministic storage and summarisation behaviour without embedding orchestration logic. This strengthens FUTON1's role as a stable substrate while avoiding scope creep into agent design or protocol negotiation.

+evidence:

- README documents the existence of the HTTP API and describes how to start it and configure its datastore.
- Focus headers and entity summaries are already emitted by the same command-service used by the API.
- XTDB-backed state is stable and consistent across CLI and API surfaces.

+next: Produce a concise API contract, document the shape of focus-header responses, and include minimal example clients demonstrating deterministic interaction.

+next-evidence:

- A single API reference file enumerating endpoints and data shapes.
- Sample clients (Clojure and one other language) communicating with the API.
- CI tests confirming that responses remain stable across versions.
- Clients that omit profile configuration correctly resolve to the documented default profile, and the '/status' surface shows the resolved datastore path.**

**! instantiated-by: Prototype 10 — World-State Access & Storage Model Documentation [Heart/已
人群中]**

+context: FUTON1 maintains the canonical world state for the stack through its combined Datascript and XTDB representation. Downstream tools—whether FUTON4, Araxa, Emacs interfaces, or ad-hoc scripts—need a clear understanding of how this state is structured and how to retrieve entities, relations, salience fields, and other stored documents without relying on internal assumptions.

+if: You want any external consumer to be able to inspect FUTON1's world state deterministically, using documented queries and data shapes rather than reverse-engineering internal code paths or storage conventions.

+however: Although the README describes persistence and hydration behaviour, there is no consolidated description of the storage model: which document types exist, how entities and relations are shaped in XTDB, how salience metadata is encoded, or which example queries reliably retrieve these structures. Export options exist, but no single reference explains their intended use.

+then: Publish a concise storage-model reference that defines FUTON1's persisted document types, key fields, and their relationships. Provide example XTDB and Datascript queries for retrieving entities, relations, salience metadata, and type information. Illustrate how downstream tools can consume this world state without depending on internal module boundaries.

+because: A clear storage-model contract makes FUTON1 a dependable foundation for the rest of the FUTON stack. It allows external tools to read and reason over the same durable world state without special-case integrations or export modes, preserving determinism and reducing the risk of accidental divergence.

+evidence:

- README documents XTDB hydration, document structure, and the relationship between XT persistence and Datascript caching.
- XT inspection workflows ('M:storage/inspect') already expose entity and relation documents.
- Demo/client workflows and ingest pipelines both write into the same store.

+next: Produce a dedicated storage-model reference; add example XTDB and Datascript queries; and clarify how exported snapshots or raw XTDB directories may be consumed by external tools.

+next-evidence:

- A single document outlining document types, field semantics, and example

queries.
- Tests confirming that example queries return stable shapes across versions.
- README or doc links showing how external tools can load and inspect the world state.

! instantiated-by: Prototype 11 – Reproducible Pilot Workflows [↗/才 ↗/程]
+context: Users may wish to construct small pilot workflows—combining ingest, demo/client interaction, and the persistent world state—to explore ideas, test downstream reasoning layers, or run controlled experiments. FUTON1 does not run these pilots itself, but provides the deterministic substrate they depend on.
+if: You want FUTON1 to support reproducible, user-defined pilot scenarios for experimental, investigative, or research-oriented purposes, without embedding pilot logic or inference in this layer.
+however: There is no standard way to package a pilot's configuration, transcripts, and expected XDB state. Any existing examples are informal or private, and users who attempt to replay scenarios must reconstruct them manually from multiple surfaces.
+then: Document how to assemble a minimal reproducible pilot using FUTON1's existing tools: define a data-root, record transcripts, capture expected outputs, and describe how to reload and inspect the resulting world state using the CLI, API, and XDB inspection functions. Provide a small example illustrating this workflow.
+because: Experimental and reasoning-oriented futons (e.g., FUTON2, FUTON4, FUTONS) require a stable, deterministic backend for storing and replaying world state. By documenting how pilots can rely on FUTON1's storage and hydration guarantees, FUTON1 supports reproducible experimentation without expanding its own scope.
+evidence:

- FUTON1 already supports deterministic transcripts and XDB-backed persistence across sessions.
- CLI and API surfaces can drive scripted workflows without modification.
- XDB inspection paths allow deterministic verification of stored state.

+next: Produce a short guide for assembling and replaying a reproducible pilot using FUTON1; include a small worked example.
+next-evidence:

- A minimal pilot example containing configuration, transcripts, and expected XDB state.
- README references describing how external futons can use such pilots as test fixtures or experimental scaffolds.

! instantiated-by: Prototype 12 – Consolidation & Packaging Pass [↗/未 ↗/甲]
+context: As FUTON1 evolves, its modules, configuration surfaces, and documentation accumulate incremental changes. Over time this can introduce small inconsistencies across apps, flags, environment variables, or test and build workflows. A periodic consolidation pass helps restore coherence without changing FUTON1's scope.
+if: You want FUTON1 to remain easy to install, run, understand, and extend by others—including tools and futons that depend on it—without requiring context from its development history.
+however: Naming and configuration conventions have drifted across modules, some flags and env vars are documented in multiple places, and not all apps follow the same patterns for logging, error handling, or test invocation. These divergences do not break functionality but reduce predictability.
+then: Perform a consolidation pass: unify naming and configuration across modules; align build and test tasks; and produce a top-level README or guide that gives a single, coherent account of FUTON1's structure, installation prerequisites, and usage patterns.
+because: A coherent packaging of FUTON1 strengthens its role as the dependable base layer for higher futons and for collaborators who rely on deterministic behaviour but may not know the internal architecture. Clear organisation reduces friction, supports reproducibility, and makes FUTON1 easier to maintain as a stable substrate.
+evidence:

- Existing READMEs document modules and workflows but vary in depth and structure.

- Test and build commands differ across 'apps/' .
- Environment-variable usage appears in multiple places with minor variations.
+next: Align naming and configuration; streamline build and test tasks; and publish a unified top-level reference document.
+next-evidence:
- Consistent flag and env-var naming across modules.
- A shared build/test pattern verified by CI.
- A single top-level README describing installation, configuration, and usage without cross-referencing scattered documents.